

Churn Dataset Analysis

- Well balanced dataset with equal distribution of classification - no need to perform sampling
- No null values present
- No Duplicated data
- Datatypes are appropriate for all features
- No correlation between features and target variable
- Data is widely distributed - has to perform normalization or standardisation for Gradient descent algorithm to reach global minima gradually.
- No linear relation among features
- No outliers present - Boxplot
- Categorical Data - Label Encoder used for 'Location' and 'Gender' features
- 'Name' and 'CustomerID' columns has no effect on target prediction - so removing both features
- Training and Testing dataset split - 70 : 30
- Data distribution is not normal or gaussian, so all features are normalised - MinMaxScaler()
- Models tested
 - Decision Tree
 - Random Forest
 - Logistic Regression
 - KNN
 - Naive Bayes
 - SVC
 - ANN
- Cross Validation - Stratifiedkfold
- Hypertunning - RandomizedSearchCV
- Feature Importance

```
Random Forest
[0.01823448 0.05012243 0.14147924 0.18567675 0.28941932 0.31506779]
['Gender' 'Location' 'Subscription_Length_Months' 'Age' 'Total_Usage_GB'
'Monthly_Bill']
XGBoost
[0.          0.01962688 0.03715258 0.14739394 0.35826558 0.43756101]
['Gender' 'Location' 'Subscription_Length_Months' 'Age' 'Total_Usage_GB'
'Monthly_Bill']
AdaBoost
[0.    0.01 0.05 0.15 0.37 0.42]
['Gender' 'Location' 'Subscription_Length_Months' 'Age' 'Total_Usage_GB'
'Monthly_Bill']
```

-
- Model Results

```

Report for logistic regression
precision    recall  f1-score   support

      0      0.51      0.69      0.58      15088
      1      0.51      0.33      0.40      14912

 accuracy      0.51      30000
  macro avg      0.51      0.51      0.49      30000
 weighted avg      0.51      0.51      0.49      30000

Report for KNN classifier
precision    recall  f1-score   support

      0      0.50      0.66      0.57      15088
      1      0.50      0.34      0.41      14912

 accuracy      0.50      30000
  macro avg      0.50      0.50      0.49      30000
 weighted avg      0.50      0.50      0.49      30000

Report for Decision Tree
precision    recall  f1-score   support

      0      0.50      0.50      0.50      15088
      1      0.49      0.50      0.50      14912

 accuracy      0.50      30000
  macro avg      0.50      0.50      0.50      30000
 weighted avg      0.50      0.50      0.50      30000

Report for Random Forest
precision    recall  f1-score   support

      0      0.50      0.54      0.52      15088
      1      0.50      0.46      0.48      14912

 accuracy      0.50      30000
  macro avg      0.50      0.50      0.50      30000
 weighted avg      0.50      0.50      0.50      30000

Report for Naive Bayes
precision    recall  f1-score   support

      0      0.51      0.69      0.58      15088
      1      0.51      0.33      0.40      14912

 accuracy      0.51      30000
  macro avg      0.51      0.51      0.49      30000
 weighted avg      0.51      0.51      0.49      30000

Report for SVM
precision    recall  f1-score   support

      0      0.50      0.60      0.55      15088
      1      0.50      0.40      0.44      14912

 accuracy      0.50      30000
  macro avg      0.50      0.50      0.50      30000

```

-
- ANN Result

```

Epoch 1/10
7000/7000 [=====] - 16s 2ms/step - loss: 0.6932 - accuracy: 0.4997
Epoch 2/10
7000/7000 [=====] - 16s 2ms/step - loss: 0.6932 - accuracy: 0.4977
Epoch 3/10
7000/7000 [=====] - 18s 3ms/step - loss: 0.6932 - accuracy: 0.4993
Epoch 4/10
7000/7000 [=====] - 14s 2ms/step - loss: 0.6932 - accuracy: 0.5010
Epoch 5/10
7000/7000 [=====] - 14s 2ms/step - loss: 0.6932 - accuracy: 0.5009
Epoch 6/10
7000/7000 [=====] - 14s 2ms/step - loss: 0.6932 - accuracy: 0.4998
Epoch 7/10
7000/7000 [=====] - 14s 2ms/step - loss: 0.6932 - accuracy: 0.4999
Epoch 8/10
7000/7000 [=====] - 14s 2ms/step - loss: 0.6932 - accuracy: 0.4981
Epoch 9/10
7000/7000 [=====] - 14s 2ms/step - loss: 0.6932 - accuracy: 0.4976
Epoch 10/10
7000/7000 [=====] - 16s 2ms/step - loss: 0.6932 - accuracy: 0.5002

```

```


# predicting the test data
y_pred = classifier.predict(xtest)
y_pred = (y_pred>0.5)
plt.plot(ann_model.history['loss'])

```

```

938/938 [=====] - 3s 3ms/step
[<matplotlib.lines.Line2D at 0x7e1e30d547f0>]

```



-
- Hypertuning for Logistic and Random Forest

```

Parameters: {'solver': 'newton-cg', 'penalty': 'l2', 'C': 0.01}
Best score is 0.5005428571428572
Parameters: {'criterion': 'gini', 'max_depth': 3, 'max_features': 2, 'min_samples_leaf': 1}
Best score is 0.5032285714285715

```

○